



# Operating Systems W3L1 - Processes & Threads: Part 1 (ctd)

▼ Class	Operating Systems
🕒 Created	@Sep 14, 2020 9:12 PM
🔗 Materials	
☑ Reviewed	<input type="checkbox"/>
▼ Type	Lecture

Regarding the *Materials*, we are continuing the notes from the last notes page. Also there was now W2L2 lecture because we didn't have two classes that week (labor day weekend).

[Operating Systems W2L1 - Processes & Threads: Part 1](#)

## A Bit About Interrupts

- All systems- big and small- have interrupts
- A **PIC** chip coordinates interruptions from processes
- Simple example of an interrupt is when you hit keys (and release keys) on your keyboard, since your OS needs to read the input
- **Interrupt:** Current running process is suspended, OS takes control

▼ When OS takes control, what modes switch?

Machine moves from **user mode** to **kernel mode**



"Kernel" is just a fancy nickname for the OS

▼ Many scenarios where an interrupt can occur (some examples)

- **Interrupt occurs due to many scenarios, for example:**
  - Time out of current running process
  - Hardware interrupt from an I/O device
  - Page fault

▼ Why might press and release be two interrupts?

1. Holding down the key is continuous
2. ???? F

▼ What is the order of events in an interrupt?

1. Hardware saves program counter, registers, etc of the current process.
2. Hardware loads program counter of the interrupt service routine (ISR).
3. ISR runs.
4. When done, the next process to run is picked.
5. Program counter, registers, etc of the picked process are loaded.
6. The picked process starts running,

- ISR is considered a *privileged process*
- If memory is not big enough, a *swap space* is used on the disk


## Moving From User to Kernel Mode

- Interrupts is an example of a process moving from user mode to kernel mode, as is a system call (which is voluntary)

# Multiprogramming

## Simple Model

- Suppose  $n$  processes in memory at once and that a process spends a fraction of  $p$  waiting for I/O
  - Probability that all processes are waiting for I/O is  $p^n$  at a given moment
  - Thus, CPU utilization is  $1 - p^n$
- ▼ Graph - CPU utilization based on a number of processes




Multiprogramming lets processes use the CPU when it would otherwise become idle.

- ▼ What don't we want CPU to be idle?


We are wasting time and resources by having processes just sitting and being idle

## Looking at OS


- ▼ Processes related to OS in different situations



(a) Separate kernel



(b) OS functions execute within user processes



(c) OS functions execute as separate processes

- (b) and (c) are used in the [literal] biggest scenarios → biggest machines use (c) due to huge number of processes running and using lots of memory
- OS processes have privilege over regular processes, but not necessarily more priorities

## Threads

▼ Multiple *threads of control* in a process, and all threads share same address space. Elaborate.

Process with four threads, the threads can share data. However, each thread needs its own program counter

- Thread of control is the execution of a program
- Simply speaking, threads are used to execute programs in parallel
- ▼ What does multithreading entail?

Process that can be broken up and operated in parallel

# Why Threads?


- Generally speaking, concurrency
  - Number of threads can change dynamically
  - Programming is "easier" with threads (and by easier, we mean not impossible because it's still very hard)
  - Benefits processes using lots of I/O and processes that overlap
- Lighter weight → faster to create and restore; much cheaper
- ▼ Two strategies for dealing with something slow
  1. Speed it up
  2. Hide it → Do something while something else is ongoing



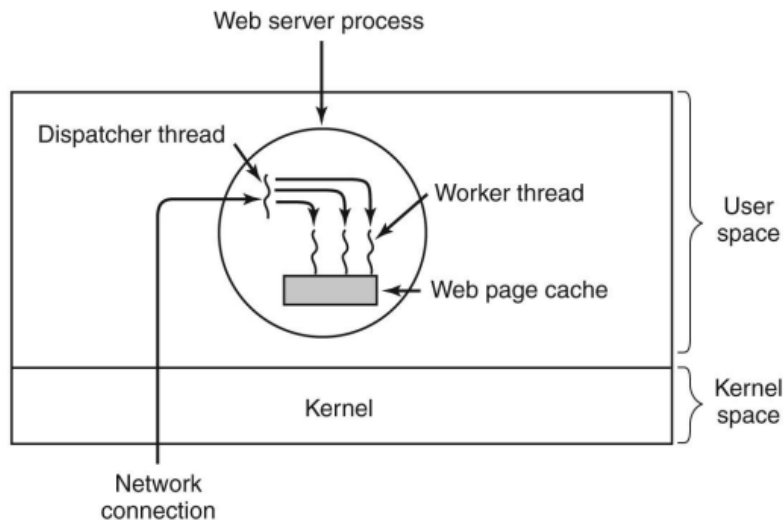
Threads share data and such, but each have their own stack

## ▼ Examples

# A Word Processor



# Multithreaded Web Server



- Even sequential processes are threaded → just one thread
  - Adding a new thread is just a unique *stack* and *program counter*
  - Threads are parts- contained by processes → **NOT a process within a process**

## Processes vs. Threads

- **Process** group resources, and **threads** are entities scheduled for execution (one of those process resources)
  - Threads sharing all fellow resources

### ▼ Why no protection between threads?

Unlike processes being protected from other processes, threads are not since threads may need to send info, share resources, etc. → they are part of the same program, so they need to work together. This process is faster without protecting threads from other threads.

- Threads can be in any of several states

### ▼ Per process vs. per thread items


Per process items	Per thread items
Address space	Program counter
Global variables	Registers
Open files	Stack
Child processes	State
Pending alarms	
Signals and signal handlers	
Accounting information	

### ▼ What is an **address space**?

You can find the [wiki page here](#), but it's- as you know from CSO- the range of discrete addresses of memory for a various number of variables and data

- Time slices allocated only to processes, not individual threads

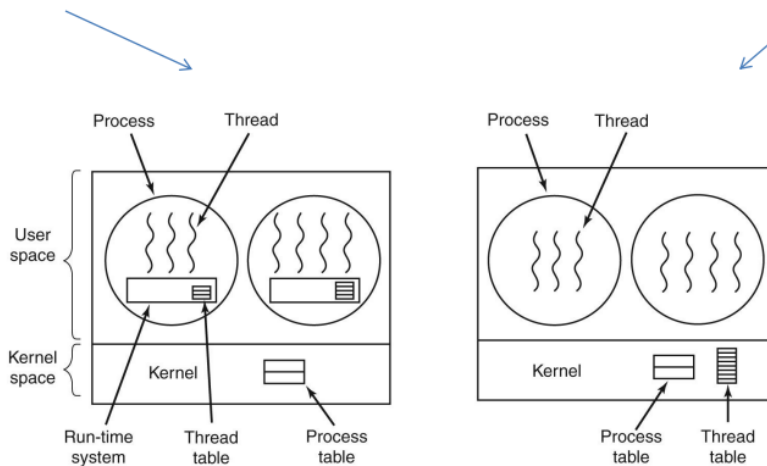
### ▼ Processes / Threads diagrams



## Where to Put The Thread Package?

User space

Kernel space



- A **thread package** is a library that helps allocate the needed "items" for threads
  - In either user space or kernel space



## Implementing Threads in User Space

- Threads are implemented by a library.
- Kernel knows nothing about threads.
- Each process needs its own private **thread table**.
- Thread table is managed by the runtime system.

## Implementing Threads in User Space

### Advantages

- Very fast thread scheduling
- Each process can have its own thread scheduling algorithm
- Scale better

### Disadvantages

- Blocking system calls can block the whole process
- Page fault blocks the whole process
- No other thread of the process will ever run unless the running thread voluntarily gives up the CPU

▼ Implementing threads in **kernel space**

## Implementing Threads in Kernel Space

- Kernel knows about and manages the threads
- No thread management is needed by the runtime.
- Creating/destroying/(other thread related operations) a thread involves a system call

## Implementing Threads in Kernel Space

### **Advantages**

- When a thread blocks (due to page fault or blocking system calls) the OS can execute another thread from the same process

### **Disadvantages**

- Cost of system call is very high

▼ What is the *runtime system*?

The environment in the OS that leaves space for programs to run