



# Operating Systems W5L2 - Race Conditions (ctd) and Deadlocks

▼ Class	Operating Systems
🕒 Created	@Sep 30, 2020 12:39 PM
📎 Materials	04 - Deadlocks.pdf
☑ Reviewed	<input type="checkbox"/>
▼ Type	Lecture

We started off by reviewing sleep/wakeup and the producer/consumer problem from last class and went over that code

[Operating Systems W5L1 - Race Conditions \(ctd\)](#)

## Solving the Lost Wakeup Problem

### ▼ What is the Lost Wakeup Problem?

When `consumer()` stops after reading `count == 0`, such as an interrupt occurring, then there's no chance to call `sleep()`. The next lecture will pick up on how we tackle this idea and (hopefully) find solutions for it.

- One solution is to add a *wakeup waiting bit*
  - An extra bit is set upon wakeup, and if the process attempts to sleep later but the bit is set, the "state" is kept awake and the bit is reset
  - Problem is... we may have more than two processes → So how many bits are we meant to use...?
- The best solution for this is **semaphores**

- Stores integer to count number of wakeups "saved for future use" — This integer is called a **semaphore**
  - This "future use" pertains to wakeup calls being saved for (and thus consumed by) processes in the future
- There are two primitives- up and down- key here at play. They are atomic operations.
  - **Up:** Increments stored integer, waking up a sleeping process (if there is one)
  - **Down:** If the integer is 0, puts process to sleep. Otherwise, it decrements and continues

#### ▼ Code using semaphores

```

#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;

/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        down(&empty);
        down(&mutex);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
}

/* TRUE is the constant 1 */
/* generate something to put in buffer */
/* decrement empty count */
/* enter critical region */
/* put new item in buffer */
/* leave critical region */
/* increment count of full slots */

void consumer(void)
{
    int item;

    while (TRUE) {
        down(&full);
        down(&mutex);
        item = remove_item();
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}

/* infinite loop */
/* decrement full count */
/* enter critical region */
/* take item from buffer */
/* leave critical region */
/* increment count of empty slots */
/* do something with the item */

```

## Mutexes

- A **mutex** is a variable that can either be locked or unlocked; in these two possible states

- It's a nickname for locks, in a way, which are used to manage critical sections
- Managed using `TSL` (see previous slides) and `XCHG` registers
- While we did say processes don't share address space, there are two ways they technically can
  1. Some of the shared data structures are stored at kernel level → Accessed via system calls
  2. Most modern operating systems offer ways for processes to "share" portions of their address space with other processes
- ▼ Message passing

This concept won't be thoroughly discussed in the course, but worth mentioning, according to Zahran.


## Forget About Sharing: How About Message Passing?

- Two primitives: send and receive
- May be used across machines
- Are system calls
  - `send(destination, &message)`
  - `receive(source, &message)`
- Issues
  - Lost acknowledgement
  - Authentication
  - performance (message passing is always slower than stuff like semaphores, ...)

---

## Deadlocks


- **Barriers** are synchronization mechanisms that are intended for a group of processes
- ▼ Barriers diagram




- We want to do everything in our power to prevent *deadlocks*

▼ What is a deadlock? What is a good analogy for it?

A deadlock is when multiple processes need to use multiple resources, but can not reach them because other processes are blocking them off. A good analogy is gridlock (as in traffic).



# Actual Deadlock



- Deadlocks occur among *processes* who need to acquire *resources* in order to *progress*.

## Resources

- A resource is anything *acquired, used, then released*.
- ▼ Preemptable vs. non-preemptable resources
  - **Preemptable:** Can be taken away from the process with no ill-effect
  - **Non-preemptable:** Cannot be taken away from the process without causing the computation to fail. In other words (since negation can be confusing), *if you take the resource away from the process, the computation (and thus the process) will fail*.
- ▼ Resource categories
  - ▼ Reusable
    - Can be safely used by only one process and is *not* depleted by that use
    - This includes processors, I/O channels, main and secondary memory, devices, and data structures (such as files, databases, and semaphores)
  - ▼ Consumable

- Can be created (produced) and destroyed (consumed)
- This includes interrupts, signals, messages, and information. Things within I/O buffers



We discussed lab1 at the end of class, approximately the last 10 minutes of the lecture recording.