



Operating Systems W10L1 - Memory Management II (ctd)

▼ Class	Operating Systems
🕒 Created	@Nov 2, 2020 6:50 PM
🔗 Materials	
☑ Reviewed	<input type="checkbox"/>
▼ Type	Lecture

Speeding Up Virtual Memory

TLB Recap

- Only a fraction of the page table is heavily used so we don't need to access the entire table. The TLB makes it so we don't have to consistently go to the page table
 - It contains page table entries for a small number of pages → A cache for the page table
 - See last lecture notes for TLB Hit and TLB Miss diagrams — Additional instructions and explanations in textbook as well
 - TLB misses are rare, and that's especially fortunate because if a miss occurs there is an additional memory cost
 - If TLB misses were frequent, we'd be even worse off in terms of performance than without it
 - TLB misses altogether occur more often than page faults
- ▼ TLB Contents

Valid	Virtual page	Modified	Protection	Page frame
1	140	1	RW	31
1	20	0	R X	38
1	130	1	RW	29
1	129	1	RW	62
1	19	0	R X	50
1	21	0	R X	45
1	860	1	RW	14
1	861	1	RW	75

- During a context switch, there are two solutions


1. Just flush out the TLB entirely
2. Add a process ID column to the TLB

The first one could be a security breach since flushing just changes valid bits to 0 and leaves the rest accessible, but the second one is larger and potentially slower

Page Table Size Problem

- We want to be sleective about what is in memory
- We do not know anything about the stack or the heap
 - Stack depends on function calls, heap depends on dynamica allocaiton

▼ Multi-Level Page Table diagram



1. One solution is to store memory address of smaller page table (start of it) in an "upper level" page table entry
 - This way we only need the small page table in memory
 - Corresponding lower-level page tables are created as needed, thus reducing memory requirement
2. Another solution is an *inverted page table* where there is one entry in the page table per page frame (as opposed to per page)
 - Saves a lot of storage, but makes virtual-to-physical storage much more difficult

Page Faults

▼ What is the definition of a page fault?

- A Page fault is when the page table is consulted, and we found that the requested page is not in the memory but in the disk.
 - That is, the page has been swapped out because the memory was full.
- Note: In reality, the OS does not wait till memory is full to start making page replacement because the OS wants to have some spare memory free just in case.